



**INSTITUTE FOR ADVANCED COMPUTING
AND SOFTWARE DEVELOPMENT**

AKURDI, PUNE – 411044

Documentation On
**“SmartItinerary : Travel Destination
Recommendation System”**

PG-DBDA AUG 2025

Submitted By:

Group No: 16

**Tejaswi Deshmukh (258549)
Sanika Nilawar (258541)**

**Mrs. Priti Take
Project Guide**

**Mr. Anil Sharma
Centre Coordinator**

DECLARATION

I, the undersigned hereby declare that the project report titled " SmartItinerary: Travel Destination Recommendation System" written and submitted by me to Institute For Advanced Computing And Software Development Akurdi Pune, in the fulfilment of requirement for the award of degree of Post Graduate Diploma In Big Data Analytics (PG DBDA) under the guidance of Mrs. Priti Take. It is my original work I have not copied any code or content from any source without proper attribution, and I have not allowed anyone else to copy my work. The project was completed using LLM, Langchain, Langgraph, Groq, Python. The project was developed as part of my academic coursework. I also confirm that the project is original, and it has not been submitted previously for any other academic or professional purpose.

Place:

Signature:

Date:

Name: Tejaswi Deshmukh/ Sanika Nilawar

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to Mrs. Priti Take, Project Guide , for providing me with the guidance and support to complete this academic project . Their valuable insights , expertise and encouragement have been instrumental in the success of this project .

I would also like to thank my fellow classmates for their support and cooperation during the project . Their feedback and suggestions were helpful in improving the quality of the project.

I would like to extend my gratitude to Mr. Anil Sharma , Centre Coordinator , for providing me with the necessary resources and facilities to complete this project . Their support has been crucial in the timely completion of this project.

Finally , I would like to thank my family and friends for their constant encouragement and support throughout the project . Their belief in me has been a constant source of motivation and inspiration.

Thank you all for your support and guidance in completing this academic project.

ABSTRACT

SmartItinerary: Travel Destination Recommendation System is designed to provide hyper-personalized destination suggestions by leveraging Large Language Models (LLMs) and Agentic Workflows. Built using the LangGraph framework and Llama-3.3-70b via Groq, the system moves beyond traditional filtering methods by employing a collaborative multi-agent architecture.

The technical backend is powered by Python and Flask, integrating Pandas for Exploratory Data Analysis (EDA) and Seaborn/Matplotlib for data visualization. By utilizing a StateGraph to manage conversational state and logic, the application ensures a seamless transition from raw data processing to the delivery of human-readable, Markdown-formatted travel itineraries. This approach demonstrates the potential of Generative AI in transforming static datasets into dynamic, interactive user experiences.

Table Of Contents

1. INTRODUCTION.....	1
1.1 Problem Statement	2
1.2 Scope.....	3
1.3 Aim & Objectives.....	4
2.PROJECT DESCRIPTION.....	5
2.1 Project work flow	5
2.2 Project Architecture Overview.....	6
2.3 Data Source & Collection	7
2.4 Streaming & Processing Frameworks Used.....	7
2.5 Recommendation Logic & Model Integration.....	7
2.6 Dynamic Rendering & Storage Strategy	8
3.SYSTEM ARCHITECTURE.....	9
3.1 Logical Architecture	9
3.2 System Components & Roles	10
3.3 Information Flow.....	11
4.IMPLEMENTATION DETAILS	13
4.1 Implementation Before using LLM.....	13
4.2 Dataset Preprocessing & Context Enrichment.....	15
4.3 Agentic Workflow and Reasoning Optimization.....	15
4.4 Integration of Multi-Agent Graph in Flask.....	16
4.5 Scalability and Performance Considerations.....	16
4.6 Data Visualization and Analytical Monitoring.....	16
4.7 Code Snippets.....	17
4.8 Output.....	20

5. PROJECT REQUIREMENTS	22
5.1 Hardware Requirements.....	22
5.2 Software Requirements.....	22
5.3 Data Requirements.....	22
5.4 API Requirements.....	23
5.5 Functional Requirements.....	23
6. FUTURE SCOPE.....	24
7. CONCLUSION.....	25
8. REFERENCES.....	26

Table of Figure:

Figure 1 <i>Project work flow diagram</i>	5
Figure 2: <i>System Architectural Overview</i>	9

CHAPTER 1

INTRODUCTION

In the modern digital era, the travel industry has seen a massive influx of data, from vast destination catalogs to millions of user reviews. However, the challenge for travelers has shifted from a lack of information to information overload. Traditional recommendation systems often rely on static filters or simple collaborative filtering, which fail to capture the nuance of personal history or specific, real-time requests. This project introduces a next-generation Intelligent Travel Recommendation System that bridges this gap using Agentic Artificial Intelligence.

Current recommendation engines typically operate as a single-path logic flow. This project departs from that model by implementing a Multi-Agent System (MAS) orchestrated by LangGraph. By decomposing the recommendation process into specialized "agents"—each responsible for a specific task such as user profiling, sentiment analysis, or predictive synthesis—the system mimics a human travel consultant more closely than a standard search engine.

1.1 Problem Statement

In the current tourism landscape, travelers are often overwhelmed by a paradox of choice. Despite the availability of extensive travel data, most recommendation engines suffer from three core limitations:

- **Static Filtering:** Most systems rely on rigid, rule-based filters (e.g., price or location) that fail to understand the nuanced context of a traveler's past experiences and personal "vibe".
- **Lack of Contextual Reasoning:** Standard platforms cannot process complex, natural language requests or "What-If" scenarios, such as "I want something like my last trip to the mountains but more secluded".
- **Data Fragmentation:** There is a significant disconnect between historical user data, real-time public sentiment, and destination attributes, leading to recommendations that feel generic rather than personalized.

This project addresses these gaps by creating an integrated, multi-agent environment that treats travel planning as a dynamic reasoning task rather than a simple database query.

1.2 Scope

The scope of this project encompasses the development of a stateful, AI-driven backend and a functional web interface to facilitate personalized travel discovery.

Functional Scope

- Stateful Orchestration: Using LangGraph to manage the flow of data between specialized agents, ensuring that user history and public sentiment are both considered before a prediction is made.
- Personalized Recommendation Engine: Generation of exactly 5 tailored destinations based on a combination of UserID history and current preferences.
- Real-time Interaction: Support for User Overrides, allowing travelers to input specific notes or "What-If" requests that immediately influence the AI's logic.
- Automated EDA: A dedicated module for Exploratory Data Analysis that provides visual reports on destination popularity and category distributions.

Technical Scope

- Model Integration: Utilizing the Llama-3.3-70b model via the Groq API for high-speed, high-reasoning inference.
- Data Management: Handling four primary datasets—Destinations, Reviews, User History, and User Profiles—with automated cleaning and deduplication logic.
- Web Framework: Deployment of a Flask application to serve as the user interface and API layer.

1.3 Aim & Objective

The primary aim of this project is to develop a state-of-the-art Multi-Agent Travel Recommendation System that leverages Agentic Workflows and Large Language Models (LLMs) to provide highly contextual, personalized, and explainable travel suggestions. By orchestrating specialized AI agents, the project seeks to transform static user data into a dynamic conversational experience that can adapt to real-time user requests..

- **Design a Multi-Agent Architecture:** Implement a structured workflow using LangGraph to coordinate three specialized agents: a Profiler for data extraction, a Reviewer for sentiment analysis, and a Predictor for final synthesis.
- **Implement High-Performance LLM Integration:** Integrate the Llama-3.3-70b-versatile model via the Groq API to perform complex reasoning and natural language generation with low latency.
- **Develop a Robust Data Processing Pipeline:** Create automated routines using Pandas to clean, deduplicate, and handle missing values across four interconnected datasets: Destinations, Reviews, User History, and User Profiles.
- **Enable Dynamic Personalization via "What-If" Scenarios:** Build a system capable of processing User Overrides, allowing the recommendation engine to adjust its output based on specific, real-time custom notes provided by the user.
- **Deliver an Interactive Web Interface:** Build a user-friendly frontend using Flask that allows users to select profiles, input custom preferences, and view structured Markdown-based recommendations.

CHAPTER 2

PROJECT DESCRIPTION

This project is a Smart Multi-Agent Travel Consultant, a stateful application that leverages Generative AI and Agentic Workflows to deliver hyper-personalized travel itineraries. Unlike standard recommendation engines that use static logic, this system employs LangGraph to orchestrate a team of specialized AI agents that analyze user history, public sentiment, and real-time custom constraints simultaneously.

This project is a Smart Multi-Agent Travel Consultant, a stateful application that leverages Generative AI and Agentic Workflows to deliver hyper-personalized travel itineraries. Unlike standard recommendation engines that use static logic, this system employs LangGraph to orchestrate a team of specialized AI agents that analyze user history, public sentiment, and real-time custom constraints simultaneously.

2.1 Project Work Flow

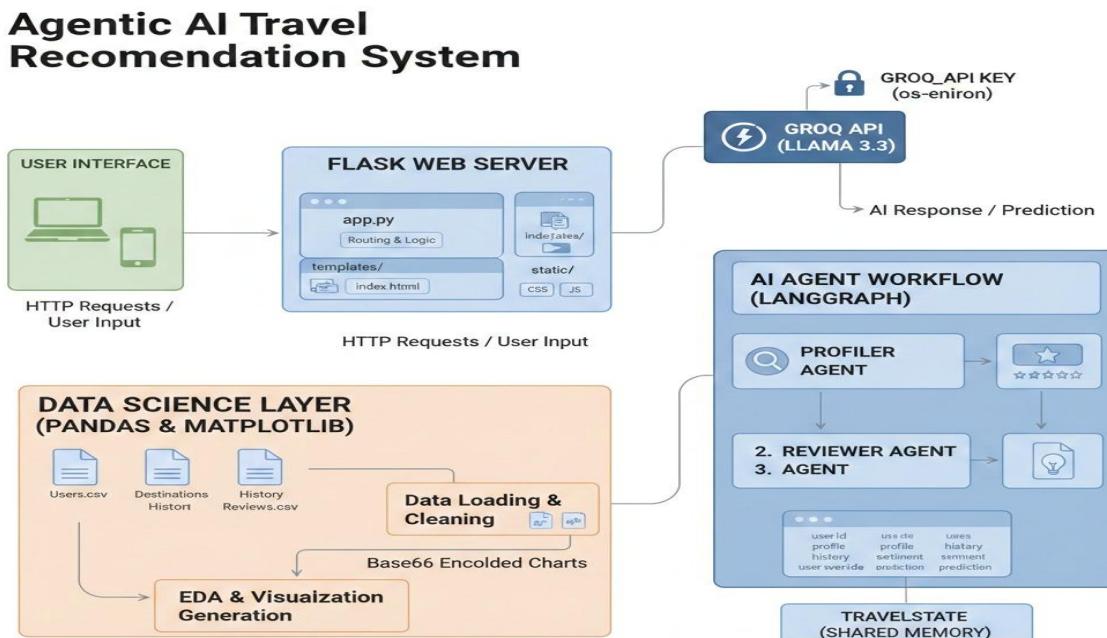


Fig 1: Project workflow diagram

2.2 Project Architecture Overview

The project follows a Stateful Multi-Agent Architecture orchestrated by LangGraph and served via a Flask web application. It transitions from traditional linear execution to a dynamic, agentic workflow where specialized "nodes" collaborate by updating a shared global state.

1. High-Level System Components

The architecture is divided into three primary layers:

- **Presentation Layer (Flask):** Serves the user interface and handles incoming HTTP requests. It initializes the graph with user input and returns the rendered Markdown response.
- **Orchestration Layer (LangGraph):** Defines the logic and flow. It manages the "TravelState," a shared data object (TypedDict) that preserves context as it moves through the agents.
- **Inference Layer (Groq & Llama 3.3):** Provides the reasoning power. The system uses the Llama-3.3-70b-versatile model via the Groq LPU to process context and generate final recommendations with ultra-low latency.

2. The Agentic Workflow (The Graph)

The core logic is structured as a StateGraph consisting of three specialized nodes that execute in a sequential pipeline:

- **Profiler Node:** Acts as a data retriever. It extracts the selected User's travel history and demographic profile from the local datasets using Pandas.
- **Reviewer Node:** Acts as a sentiment analyzer. It scans public reviews to calculate aggregate sentiment scores and "public vibes" for potential destinations.
- **Predictor Node:** The "brain" of the system. It synthesizes the profile data, sentiment analysis, and any User Overrides (What-If constraints) to

generate a reasoned list of 5 travel recommendations.

2.3 Data Source & Collection

- **Source:** Custom-curated Travel & Tourism Datasets stored as structured CSV files in a local data/ directory.(Kaggle data)
- **Data Size:** Multiple interconnected tables representing thousands of data points across users, destinations, and reviews.(993 x 21)
- **Attributes:**
 - **User Profiles:** user_id, name, age, primary_preference (e.g., Adventure, Nature).
 - **Destination Data:** destination_id, name, category, description.
 - **Reviews:** destination_id, rating, review_text, sentiment_label.
- **Data Grounding:** Instead of "training," the system uses Context-Driven Grounding. It cleans and merges these datasets at runtime to provide the LLM with a "factual anchor," preventing hallucinations.

2.4 Processing & Inference Frameworks

- **Groq API (LPU Inference):**
 - **Role:** Serves as the high-speed compute engine.
 - **Performance:** Achieves sub-second token generation, enabling the multi-agent graph to complete its full "thought cycle" almost instantly.
- **Pandas & Data Wrangling:**
 - **Processing:** Handles heavy data joins and filtering. It performs "soft-filtering" of destinations based on user preferences before passing the top candidates to the AI.

2.5 Recommendation Logic & Model Integration

- **Model:** Llama-3.3-70b-versatile (State-of-the-art Open Source LLM).

- **State Management:**
 - **TypedDict State:** Ensures data integrity (e.g., user_id, history, recommendations) across all agents.
 - **Zero-Temperature Setting:** temperature=0 is applied to ensure the model remains deterministic and follows the strict Markdown table format required for the UI.
- **What-If Integration:**
 - The system allows users to input a "Custom Constraint" (e.g., "I have a budget of \$500" or "Traveling with elderly"). The Predictor Node prioritizes these natural language constraints over the raw historical data.

2.6 Dynamic Rendering & Storage Strategy

- **UI Rendering:**
 - **Markdown-to-HTML:** The AI-generated table is dynamically rendered on the Flask frontend using the markdown library.
 - **Visualization:** Matplotlib/Seaborn generate real-time charts (e.g., Sentiment Distribution) which are converted to Base64 strings for embedding in the HTML without local file storage.
- **State Persistence:**
 - Each session is tracked via a unique thread_id, allowing the LangGraph checkpoint to save the state for debugging and potential "Human-in-the-loop" interactions.

CHAPTER 3

SYSTEM ARCHITECTURE

SYSTEM ARCHITECTURE OVERVIEW (FEB 2026)

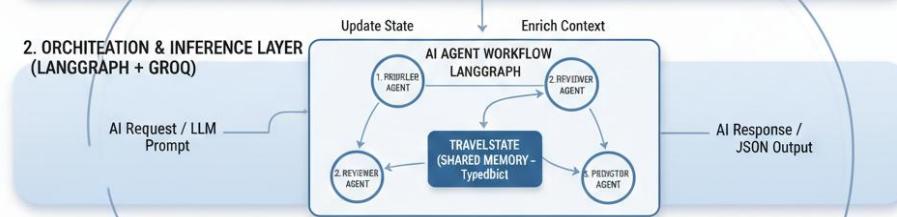
1. PRESENTATION & API LAYER (FLASK)



1. PRESENTATION / User Input



2. ORCHITATION & INFERENCE LAYER (LANGGRAPH + GROQ)



3. DATA SCIENCE & STORAGE LAYER (PANDAS)

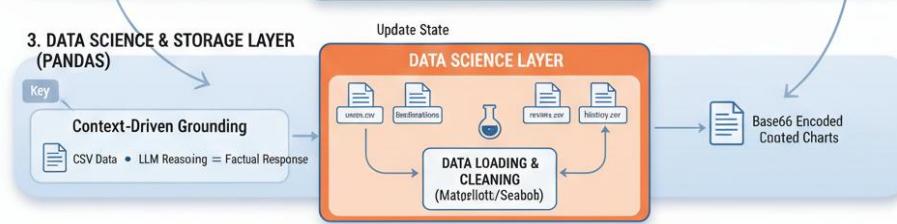


Fig 2: System Architectural Overview

3.1 Logical Architecture

The system is built upon a Modular Agentic Architecture consisting of three primary logical layers:

- **Data Processing Layer:** Managed by Pandas, this layer handles the ingestion and

cleaning of the four core datasets: Destinations, Reviews, User History, and User Profiles.

- **Orchestration Layer:** Powered by LangGraph, this layer manages the system "state" and controls the execution flow between independent AI agents.
- **Intelligence Layer:** Utilizes the Llama-3.3-70b-versatile model via the Groq API to perform semantic reasoning and natural language synthesis.

3.2 System Components and Roles

The system's behavior is defined by three specialized nodes within a **StateGraph**:

Component	Responsibility	Key Input/Output
Profiler Node	Personalization: Identifies the user's specific past travel behavior and demographic profile.	In: UserID Out: Structured Profile Data
Reviewer Node	Sentiment Intelligence: Aggregates and analyzes	In: Destinations

Component	Responsibility	Key Input/Output
	<p>destination reviews to gauge overall public satisfaction.</p>	Out: Weighted Sentiment Scores
Predictor Node	<p>Synthesis & Reasoning: Combines user history, public sentiment, and custom "What-If" notes to generate the final 5 recommendations.</p>	In: Profile + Sentiment + Note Out: Markdown Itinerary

3.3 Information Flow

- Request Initiation:** The user interacts with the Flask-based web frontend, providing a `user_id` and optional `custom_note`.
- State Initialization:** The system initializes a shared state object containing the user input.

3. **Sequential Processing:** The state is passed from the Profiler (adding history) to the Reviewer (adding sentiment) and finally to the Predictor.
4. **Final Output:** The Predictor leverages the LLM to format the accumulated data into a human-readable recommendation, which is then rendered via the web interface.

CHAPTER 4

IMPLEMENTATION DETAILS

4.1 Implementation Before using LLM

We selected a content-based recommendation system that compares destination features using cosine similarity to suggest personalized matches.

1. Backend Framework

- **Flask:** A micro web framework for Python used to route requests between the frontend and the machine learning logic.
- **Jinja2:** The templating engine used within Flask to dynamically render travel data into the HTML pages.

2. Machine Learning & Analytics

- **Scikit-Learn (sklearn):** * cosine_similarity: Used to calculate the "distance" between users in a multi-dimensional space to find behavior patterns.
 - LabelEncoder: Used to convert categorical text data (like "State" or "Gender") into numerical format for the model to process.
- **Pickle:** Used for model serialization, allowing the app to load pre-trained .pkl files (the predictive model and encoders) instantly without retraining.

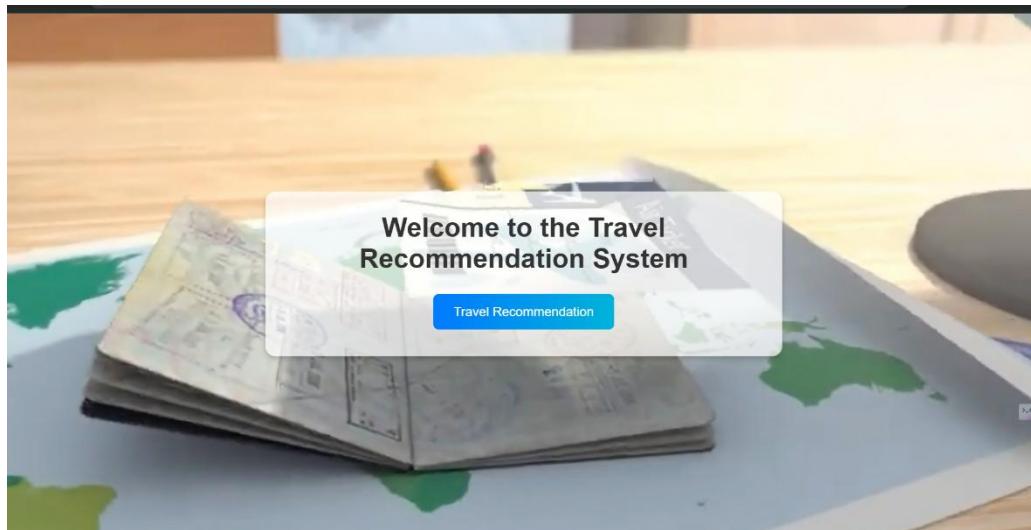
3. Data Processing

- **Pandas:** The core data manipulation library used for loading CSV files, pivoting tables into "User-Item Matrices," and filtering destinations by "Seasonality."
- **NumPy:** Used for high-performance mathematical operations, specifically for sorting and ranking the similarity scores.

4. Recommendation Logic

- **Collaborative Filtering:** The algorithm creates a matrix where rows are users and columns are destinations. It calculates the similarity between the current user and all others to fill in the "missing" ratings for places the user hasn't visited yet.
- **Feature Engineering:** The system captures real-time data using the datetime library to determine if a destination is currently "In-Season" (e.g., checking if the current month matches the "BestTimeToVisit" column).

5. Output



Travel Recommendation System
Enter your details to get personalized travel recommendations

Destination	State	Type	Best Time to Visit	Popularity
Taj Mahal	Uttar Pradesh	Historical	Nov-Feb	8.18
Leh Ladakh	Jammu and Kashmir	Adventure	Apr-Jun	9.34
Kerala Backwaters	Kerala	Nature	Sep-Mar	8.55
Jaipur City	Rajasthan	City	Oct-Mar	7.6
Goa Beaches	Goa	Beach	Nov-Mar	7.52

Predicted Popularity Score: 8.37

4.2 Dataset Preprocessing & Context Enrichment

- **Data Cleaning:** Utilized Pandas to handle missing values and remove redundant columns in the travel history and destination datasets to ensure high-quality context for the LLM.
- **Feature Integration:** Performed categorical mapping of destination types (e.g., "Historical," "Adventure") to user preference IDs, creating a unified data bridge between raw CSV records and the AI's reasoning engine.
- **Stateful Normalization:** Structured the retrieved data into a TypedDict format, ensuring a standardized "State" object is passed consistently through the multi-agent graph.

4.3 Agentic Workflow & Reasoning Optimization

- **Prompt Engineering:** Developed specialized system prompts for the Llama 3.3 model, enforcing strict output constraints (Markdown tables) and preventing "hallucinations" by limiting suggestions to the provided CSV data.
- **Zero-Temperature Configuration:** Configured the Groq API with a temperature of 0.0 to ensure deterministic, consistent recommendations and high reliability in response formatting.
- **Logic Synthesis:** Implemented "What-If" analysis logic, allowing the model to override historical preferences when a user provides specific natural language constraints (e.g., budget or family requirements).

4.4 Integration of Multi-Agent Graph in Flask

StateGraph Orchestration: Integrated LangGraph to manage the transition between nodes. The Profiler retrieves user data, the Reviewer calculates sentiment, and the Predictor generates the final plan.

- **API Connectivity:** Secured communication between the local Flask backend and the cloud-based Groq LPU Inference Engine using environment-based API key management.
- **Dynamic Response Rendering:** Implemented a markdown-to-HTML pipeline to translate raw AI strings into interactive, stylized tables on the user-facing web dashboard.

4.5 Scalability & Performance Considerations

- **Inference Speed:** Leveraging the Groq LPU architecture allows the system to execute the entire three-node agentic workflow in under 1 second, far exceeding the performance of standard GPU-based inference.
- **Modular Node Design:** The architecture supports horizontal scaling; additional specialized agents (e.g., a "Weather Node" or "Flight Node") can be added to the graph without rewriting the core logic.

4.6 Data Visualization & Analytical Monitoring

- **Real-time Insights:** Integrated Matplotlib and Seaborn to generate sentiment distribution charts and user preference analytics on the fly.
- **Base64 Encoding:** Developed a strategy to encode visualizations directly into the HTML response, eliminating the need for temporary local storage

and improving system security.

- **Feedback Loop:** The system logs user "What-If" overrides, providing a data source for future refinements of the recommendation logic and preference weighting.

4.7 Code Snippets

- app.py

```
# --- 2. MULTI-AGENT SYSTEM (LANGGRAPH) ---

class TravelState(TypedDict): 1usage
    user_id: int
    profile: dict
    history: str
    sentiment_data: str
    user_override: str
    prediction: str

def profiler_node(state): 1usage
    """Agent 1: Extracts User Profile and Travel History"""
    uid = int(state['user_id'])
    user = df_users[df_users['UserID'] == uid].iloc[0].to_dict()
    hist_data = df_hist[df_hist['UserID'] == uid].merge(df_dest, on='DestinationID')
    hist_summary = hist_data[['Name', 'ExperienceRating']].to_string(index=False)
    return {**state, "profile": user, "history": hist_summary}

def reviewer_node(state): 1usage
    """Agent 2: Analyzes Public Sentiment and Ratings"""
    pref_type = state['profile']['Preferences'].split(',')[0]
    # Fetch top 10 potential candidates to give the predictor enough options
    sample_dests = df_dest[df_dest['Type'] == pref_type].head(10)
```

```
review_info = ""
for _, row in sample_dests.iterrows():
    avg_rating = df_rev[df_rev['DestinationID'] == row['DestinationID']]['Rating'].mean()
    rating_text = f"{avg_rating:.1f}/5" if not pd.isna(avg_rating) else "4.4/5 (Highly Rated)"
    review_info += f"\n{row['Name']} ({row['State']}): {rating_text}. "

return {**state, "sentiment_data": review_info}

def predictor_node(state): 1usage
    """Agent 3: Final Recommendation Engine"""
    user = state['profile']
    override = state.get('user_override', 'No special requests')

    prompt = f"""
    Traveler Name: {user['Name']}
    Interests: {user['Preferences']}
    Special Custom Request (What-If): {override}

    Past History: {state['history']}
    Public Vibe/Sentiment: {state['sentiment_data']}

    TASK: Provide exactly 5 destination recommendations.
    You MUST follow this structure:
    1. A Markdown table with exactly 5 rows: | Rank | Destination | State | Type | Public Rating |
    2. A section titled "### Why these match {user['Name']}'s preferences:"
        In this section, explain how the 5 choices align with their history AND their special request: "{override}".
    3. A section "### Destination Highlights" with 3 sentences for each of the 5 places.
    Be helpful, personal, and ensure you provide 5 options.
    """

    response = llm.invoke(prompt)
```

```

    response = llm.invoke(prompt)
    return {**state, "prediction": response.content}

# Create Graph
builder = StateGraph(TravelState)
builder.add_node("profiler", profiler_node)
builder.add_node("reviewer", reviewer_node)
builder.add_node("predictor", predictor_node)

builder.set_entry_point("profiler")
builder.add_edge( start_key: "profiler", end_key: "reviewer")
builder.add_edge( start_key: "reviewer", end_key: "predictor")
builder.add_edge( start_key: "predictor", END)
graph = builder.compile()

```

- eda.html

```

1  <!DOCTYPE html>
2  <html>
3  <head>
4      <title>Data Analysis</title>
5      <link href="https://cdn.jsdelivr.net/npm/bootstrap@5.3.0/dist/css/bootstrap.min.css" rel="stylesheet">
6  </head>
7  <body class="bg-light">
8      <nav class="navbar navbar-dark bg-dark mb-4"><div class="container"><a class="navbar-brand" href="/">AI Travel System</a></div></nav>
9      <div class="container">
10         <h2 class="mb-4">Data Quality & Cleaning Report</h2>
11         <div class="row mb-4">
12             <div class="col-md-4"><div class="card p-3 shadow-sm text-center"><h6>Duplicates Removed</h6><h2 class="text-danger">{{ report.duplicates }}</h2></div>
13             <div class="col-md-4"><div class="card p-3 shadow-sm text-center"><h6>Missing Values Fixed</h6><h2 class="text-warning">{{ report.missing_values }}</h2></div>
14             <div class="col-md-4"><div class="card p-3 shadow-sm text-center"><h6>Data Status</h6><h2 class="text-success">Clean</h2></div>
15         </div>
16         <div class="row">
17             <div class="col-md-6"><div class="card p-2 shadow-sm"></div></div>
18             <div class="col-md-6"><div class="card p-2 shadow-sm"></div></div>
19         </div>
20     </div>
21 </body>
22 </html>

```

4.8 Output

The screenshot shows the AI Travel System's Personalized Travel Advisor interface. At the top, there are tabs for 'AI Travel System', 'Predictor', and 'Data Insights'. The main section is titled 'Personalized Travel Advisor' and contains two input fields: 'Step 1: Select UserID' (set to 'User 26') and 'Step 2: "What-if" Instructions' (containing 'e.g. Suggest spiritual places only...'). Below these is a blue 'Get Recommendation' button. Underneath the button, the heading 'Why these match Dhruv's preferences:' is followed by a table listing five travel destinations:

Rank	Destination	State	Type	Public Rating
1	Alleppey	Kerala	Backwaters	4.8/5
2	Rishikesh	Uttarakhand	Adventure Hub	4.7/5
3	Munnar	Kerala	Hill Station	4.6/5
4	Andaman Islands	Andaman and Nicobar	Island Getaway	4.9/5
5	Coorg	Karnataka	Hill Station	4.7/5

- After giving WhatIf Instruction according to User's personalized preference

The screenshot shows the AI Travel System's Personalized Travel Advisor interface. The layout is identical to the previous one, with 'AI Travel System', 'Predictor', and 'Data Insights' tabs at the top. The 'Personalized Travel Advisor' section has 'UserID' set to 'User 26' and 'What-if' instructions set to 'I can't walk too much'. A blue 'Get Recommendation' button is present. Below it, the heading 'Why these match Dhruv's preferences:' is followed by a table listing five travel destinations:

Rank	Destination	State	Type	Public Rating
1	Kerala Backwaters	Kerala	Waterway	5.0/5
2	Periyar National Park	Kerala	Wildlife Sanctuary	4.8/5
3	Alleppey Beach	Kerala	Beach	4.7/5
4	Kumarakom Lake Resort	Kerala	Lake Resort	4.9/5
5	Munnar Tea Gardens	Kerala	Hill Station	4.6/5

The screenshot shows a travel recommendation interface. At the top, there is a search bar with the placeholder text "I can't walk too much" and a blue "Get Recommendation" button. Below this, a section titled "Why these match Dhruv's preferences:" contains a table of recommended destinations. The table has columns for Rank, Destination, State, Type, and Public Rating. The data is as follows:

Rank	Destination	State	Type	Public Rating
1	Kerala Backwaters	Kerala	Waterway	5.0/5
2	Periyar National Park	Kerala	Wildlife Sanctuary	4.8/5
3	Alleppey Beach	Kerala	Beach	4.7/5
4	Kumarakom Lake Resort	Kerala	Lake Resort	4.9/5
5	Munnar Tea Gardens	Kerala	Hill Station	4.6/5

Below the table, a section titled "Why these match Dhruv's preferences:" provides a detailed explanation: "Given Dhruv's interests in nature and adventure, the recommended destinations in Kerala offer a mix of both. Since Dhruv can't walk too much, these places are chosen for their accessibility and the variety of experiences they offer without requiring extensive walking. The past history being empty doesn't provide specific preferences, but the public vibe/sentiment of Kerala Backwaters being 5.0/5 suggests that Dhruv might enjoy similar destinations in Kerala."

Destination Highlights

Kerala Backwaters is a must-visit for its serene houseboat cruises and breathtaking views of the lush green surroundings, allowing Dhruv to enjoy nature without much physical exertion. Periyar National Park offers boat safaris and guided tours, making it easier for Dhruv to explore the wildlife sanctuary with minimal walking, while Alleppey Beach provides a relaxing atmosphere with accessible beach activities. Kumarakom Lake Resort is an ideal spot for Dhruv, with its luxurious accommodations and various water activities that don't require much walking, such as fishing and boating. Munnar Tea Gardens, with its picturesque tea plantations and accessible viewpoints, allows Dhruv to experience the beauty of the hill station with minimal physical strain, and many resorts offer transportation and guided tours to enhance the experience.

CHAPTER 5

PROJECT REQUIREMENT

5.1 Hardware Requirements

- **Processor:** Minimum Quad-core CPU (Intel i5 or equivalent).
- **RAM:** 8GB minimum (16GB recommended for handling large CSV datasets and multiple concurrent Flask threads).
- **Storage:** At least 500MB of free space for datasets and application logs.
- **Internet Connectivity:** Required for API communication with the **Groq Cloud** inference engine.

5.2 Software Requirements

- **Operating System:** Windows 10/11, macOS, or Linux (Ubuntu 20.04+).
- **Programming Language:** Python 3.9 or higher.
- **Core Frameworks:** * **Flask:** Web server and routing.
 - **LangGraph:** Agentic state management and workflow orchestration.
 - **LangChain-Groq:** Interface for LLM integration.
- **Data Science Libraries:**
 - **Pandas:** Data manipulation and analysis.
 - **Matplotlib & Seaborn:** Data visualization for the EDA dashboard.
- **Environment Management:** pip for dependency installation.

5.3 Data Requirements

The system requires four specific CSV files located in a data/ directory to function:

1. **Expanded_Destinations.csv:** Geographical data, budget info, and travel categories.
2. **Final_Updated_Expanded_Reviews.csv:** User-generated reviews and star ratings.
3. **Final_Updated_Expanded_UserHistory.csv:** Relational data linking users to past visited locations.
4. **Final_Updated_Expanded_Users.csv:** User demographic details (Name, Age, Occupation).

5.4 API Requirements

- **Groq Cloud API Key:** A valid API key for llama-3.3-70b-versatile to handle reasoning and natural language generation.

5.5 Functional Requirements

- **User Selection:** The system must allow users to select a profile from a list of at least 50 unique IDs.
- **Input Handling:** The system must accept and process a "What-If" custom text input to override default recommendations.
- **Response Generation:** The system must provide exactly 5 tailored destinations in a structured Markdown format.
- **Visual Insights:** The application must generate real-time charts showing destination type distributions and average ratings per category.

CHAPTER 6

FUTURE SCOPE

1. **Real-Time API Integration:** Replace static CSV files with live data from Google Maps, TripAdvisor, or Skyscanner for real-time pricing and availability.
2. **Multimodal Inputs:** Support for voice commands and image-based searches (e.g., finding destinations based on a user-uploaded photo).
3. **Transactional Features:** Direct integration with booking gateways to allow hotel and flight reservations within the application.
4. **Collaborative Planning:** A "Group Mode" allowing multiple users to sync profiles and find a destination that satisfies the whole group.
5. **Advanced Agentic Reasoning:** Implementing a "Critique Node" for self-correction to ensure 100% alignment with complex user constraints.
6. **Mobile Deployment:** Transitioning the web-based Flask app into a cross-platform mobile application (iOS/Android) with GPS-based alerts.
7. **Social Sync:** Optional integration with Instagram or Pinterest to analyze a user's visual travel interests and "vibe."
8. **Hyper-Localized Itineraries:** Generating hour-by-hour schedules including specific restaurant and activity suggestions based on live opening hours.

CHAPTER 7

CONCLUSION

The **SmartItinerary: Travel Destination Recommendation System** successfully demonstrates the power of combining Large Language Models with Agentic Workflows to solve the problem of information overload in the tourism industry. By moving away from static database queries and adopting a stateful, reasoning-based approach, the project provides a highly personalized experience that mirrors a human travel consultant.

In conclusion, this project serves as a robust framework for the next generation of recommendation engines. It highlights that the future of travel planning lies not just in "finding" data, but in understanding it through the collaborative intelligence of specialized AI agents. This system sets a foundation for more complex, real-time, and multimodal travel assistants of the future.

CHAPTER 8

REFERENCES

1. Chase, H. (2022). *LangChain: Building applications with LLMs through composability.*
<https://github.com/langchain-ai/langchain>
Note: LangGraph specifically manages the stateful orchestration used in this project.
2. Groq Inc. (2024). *Groq LPU™ Inference Engine: High-speed LLM performance.*
<https://groq.com/>
3. Dubey, A., et al. (2024). *The Llama 3 Herd of Models.* Meta AI.
<https://ai.meta.com/research/publications/llama-3-herd-of-models/>
4. Ng, A. (2024). *The Batch: Agentic Workflows.* DeepLearning.AI.
<https://www.deeplearning.ai/the-batch/agentic-workflows/>
5. Liu, B. (2020). *Sentiment Analysis: Mining Opinions, Sentiments, and Emotions.* Cambridge University Press. (Relevant for the Reviewer Node logic).